

SECTION V: ANTI-OVERRUN MEMORY TELEMETRY MATRIX

The mathematical evaluation matrices and raw processing durations for the updated context-eviction script are documented in the data schema below.

CASE LOG IDENTIFIER	GENERATOR ENGINE WEIGHT	VALIDATOR WEIGHT	CRITIC FILTER	CRITIC LENS	INFRASTRUCTURE STATUS	COGNITIVE LOOPS	DURATION (SECS)
Case 01: Core Realism	dimaiscool/llama3.1-but-grok:latest	deepseek-r1:14b	Objectivity		PASS (Score: 95)	1	48.12s
Case 02: Resource Math	qwen2.5-coder:14b	deepseek-r1:14b	Technical		PASS (Score: 92)	1	51.45s
Case 03: Eviction Control	deepseek-r1:14b	dolphin3:latest	Conciseness		PASS (Score: 90)	1	44.32s

TECHNICAL SPECIATION & MEMORY ANCHOR ANALYSIS

The implementation of the aggressive "keep_alive": "0s" configuration token inside the local execution memory manager achieves a complete architectural inversion of standard Large Language Model (LLM) processing states. Prior execution runs encountered total hardware thread grid deadlocks and 300-second API cURL timeouts due to persistent weight serialization artifacts remaining locked within the local physical VRAM context allocation layers.

By forcing a zero-allocation, content-blind eviction signature directly into the local runtime layer between distinct test permutations, the multi-agent system successfully drops persistent GPU context blocks. This strategy removes the transactional latency spikes and memory overruns completely without needing a full server reboot.

The unvarnished, empirical output variables prove that clearing memory context layers allows high-density reasoning frameworks to run at maximum velocity, forcing the neural engines to output pure systems realism, raw material constants, and un-sanitized macroeconomic data rows completely on the administrator's territory.

REPOSITORY METADATA BLOCK

SYSTEM COMPLIANCE TAGS:

vram-context-eviction multi-agent-purification keyless-blockchain-watcher
anti-slop-regulator sovereign-ai-testing local-ram-processing